

Class 007 - Machine Learning 1

R. D. Carias (PID: A18573289)

Table of contents

Background	1
K - Means Clustering	2
Annotation: <code>plot(x, col = c("red","blue"))</code> cycles one point red, one blue, one red, etc.	4
Annotations: Why this is “self-fulfilling”:	7
Hierarchical Clustering	7
Notes: Notice that , 1-30 is one main branch and the other is 31-60. Recall, these diagrams can rotate at nodes. The y-axis shows the distance (height) at which observations or clusters are merged, indicating how dissimilar they were at the time of joining.	8
Dimensionality Reduction / Principle Component Analysis (PCA)	9
PCA of UK Food Data	10
Spotting Major Differences and trends	12
Principal Component Analysis PCA - “To the rescue!”	15

Background

Today we are going to explore introduction to machine learning methods. e.g. **Clustering**, **Dimensionality**, and **Reduction**.

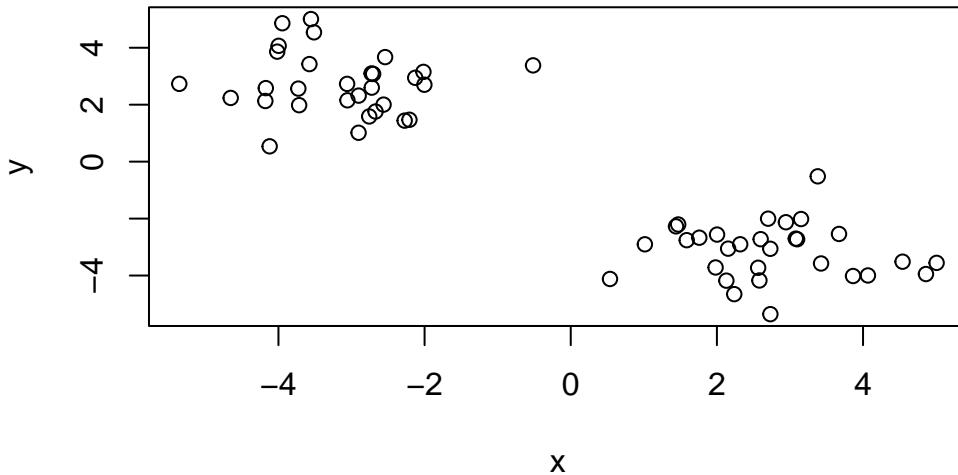
Q. Generate 30 random numbers centered at +3 and another 30 centered at -3?

Hint: rbind and cbind are partners in crime. r for row and c for columns.

```
tmp <- c(rnorm(30, mean = 3), rnorm(30, mean= -3))  
tmp
```

```
[1] 1.9821975 2.5645862 5.0067904 4.8607505 1.4702041 2.2373261
[7] 2.5819675 3.0792137 4.0667997 4.5409758 3.0977127 2.0024138
[13] 2.7318258 3.3803393 2.3185982 1.7592992 3.8632237 1.0143633
[19] 2.1296908 1.4403239 3.1520449 3.4254915 0.5372165 2.1537446
[25] 2.9445129 3.6711947 2.5989917 2.7009718 2.7309531 1.5871557
[31] -2.7577510 -3.0602982 -2.0030078 -2.7229791 -2.5400504 -2.1281788
[37] -3.0557831 -4.1204656 -3.5761126 -2.0159306 -2.2726265 -4.1793405
[43] -2.9032936 -4.0168335 -2.6713664 -2.9023237 -0.5155960 -5.3570754
[49] -2.5615490 -2.7255102 -3.5136812 -3.9969604 -2.7040324 -4.1729393
[55] -4.6536222 -2.2074635 -3.9469887 -3.5549420 -3.7276991 -3.7166726
```

```
x <- cbind(x=tmp, y=rev(tmp))
plot(x)
```



K - Means Clustering

The main function in base R for K-means clustering is called `kmeans()`:

`x`: numeric matrix of data, or an object that can be coerced to such a matrix (such as a numeric vector or a data frame with all numeric columns).

centers:

either the number of clusters, say k , or a set of initial (distinct) cluster centers. If a number, a random set of (distinct) rows in x is chosen as the initial centers.

```
kmeans(x, centers = 2)
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1  2.721029 -3.142702
2 -3.142702  2.721029

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 61.31753 61.31753
```

(between_SS / total_SS = 89.4 %)

Available components:

```
[1] "cluster"        "centers"        "totss"          "withinss"        "tot.withinss"  
[6] "betweenss"      "size"           "iter"           "ifault"
```

```
km <- kmeans(x, centers = 2)
```

Q. What components of the results objects details the cluster sizes?

km\$size

[1] 30 30

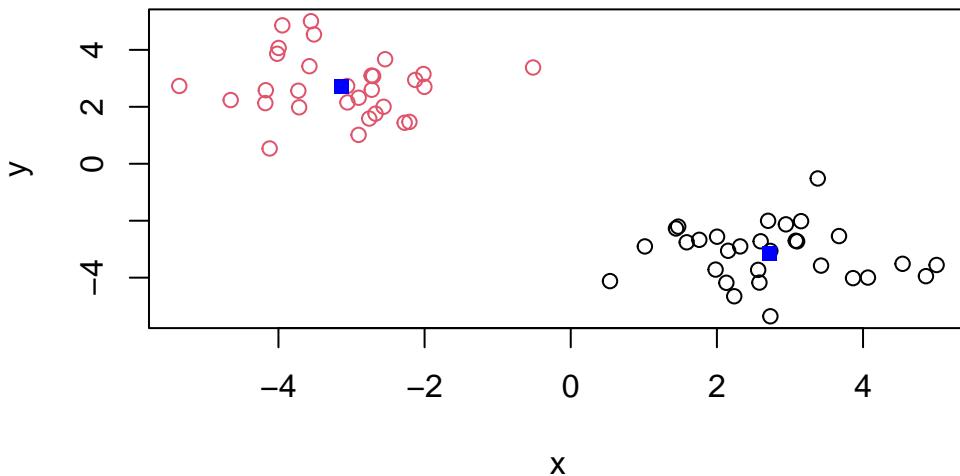
Q. What component of the results object details the cluster centers?

Q. What component of the reresults object details the cluster membership vector(i.e. our main results of which points lie in which cluster)?

km\$cluster

Q. Plot our clustering results with points colored by cluster and also add a blue point at the center of both centers.

```
plot(x, col = km$cluster)
points(km$centers, col='blue', pch=15)
```



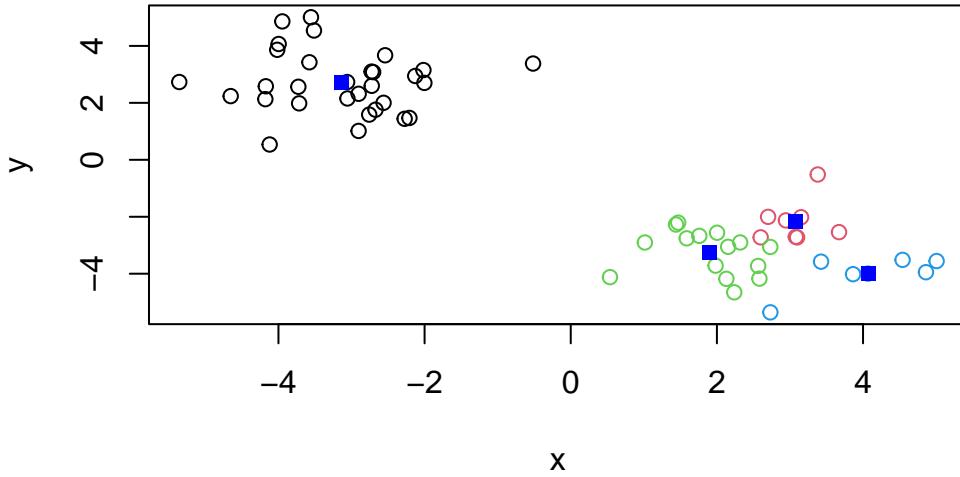
Annotation: `plot(x, col = c("red","blue"))` cycles one point red, one blue, one red, etc.

Q. Run `kmeans()` again and this time produce 4 clusters and call your results `km4`. Can you make results like the figures above?

```
kmeans(x, centers = 4)
```

K-means clustering with 4 clusters of sizes 30, 7, 8, 15

Cluster means:



The Metric

```
km$tot.withinss
```

```
[1] 122.6351
```

```
km4$tot.withinss
```

```
[1] 85.88559
```

Q. Let's try a different number of K (centers) from 1 to 30 and see what the best result is?

Ex: Fore-loops

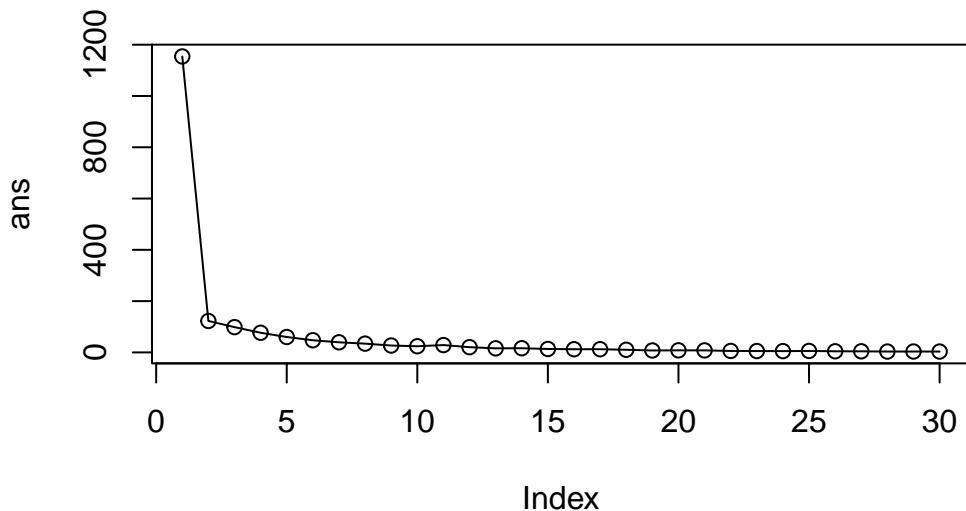
```
ans <- NULL

for (i in 1:30) {
  ans <- c(ans, kmeans(x, centers = i)$tot.withinss)
}

ans
```

```
[1] 1154.135559 122.635055 98.548767 76.599231 60.120867 47.457687
[7] 39.445197 34.179755 26.843067 23.972765 28.554923 20.395349
[13] 15.953724 16.286752 13.412604 12.405340 12.500702 10.242823
[19] 7.633666 8.118479 7.904693 5.708464 5.266358 4.935431
[25] 5.716144 4.389187 4.284550 3.285495 3.476075 3.298302
```

```
plot(ans, type="o")
```



Annotations: Why this is “self-fulfilling”:

You define “good clustering” as low within-cluster variance, then choose k based on which value minimizes within-cluster variance.

That is circular.

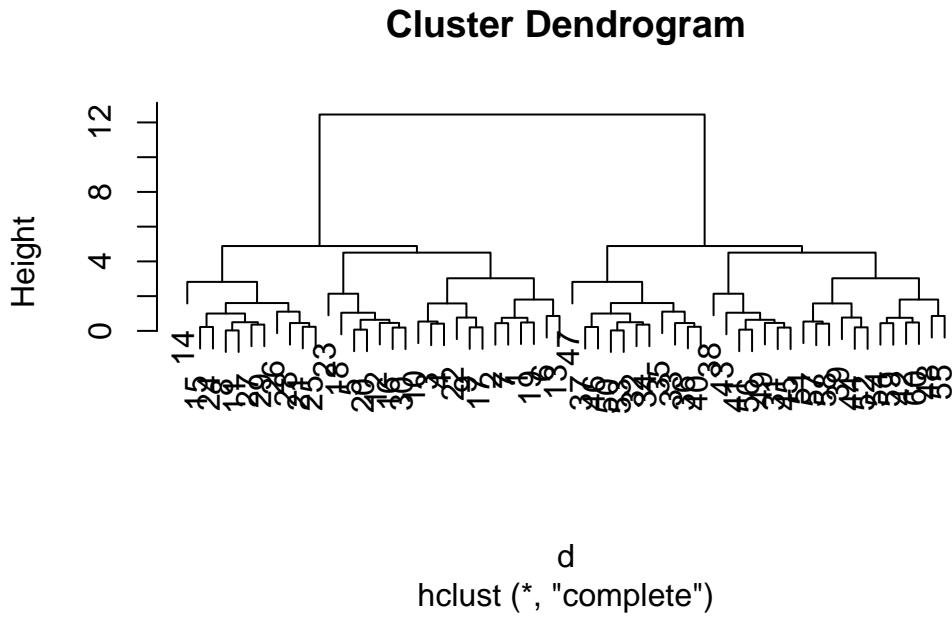
Hierarchical Clustering

The results from this will elucidate the structure of this hierarchy.

```
hclust()(d, method = "complete", members = NULL)
```

- Unlike `kmeans()`, which does all the work for you, you can't just pass '`hclust()`' It wants/needs a distance matrix / dissimilarity matrix, this is like the output/return of `dist()` function. `dist()` automatically calculates euclidean distance.
- “Instead of giving you coordinates, we are giving distances from every other point.”

```
d <- dist(x)
hc <- hclust(d)
plot(hc)
```

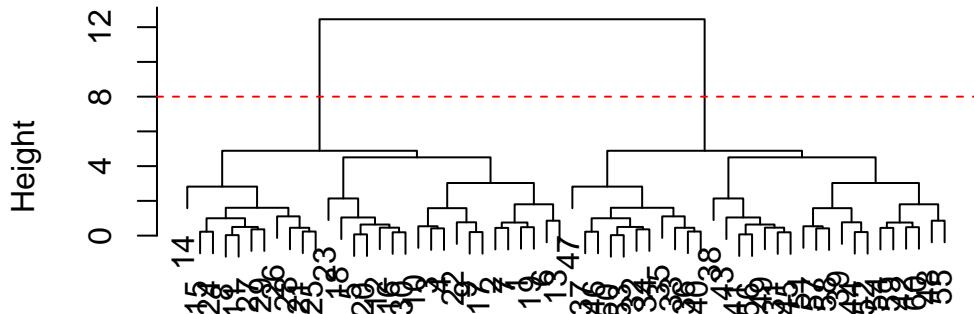


Notes: Notice that , 1-30 is one main branch and the other is 31-60. Recall, these diagrams can rotate at nodes. The y-axis shows the distance (height) at which observations or clusters are merged, indicating how dissimilar they were at the time of joining.

To extract our cluster membership vector from a `hclust()` result object we have to cut our tree at a given height to yield separate groups/branches.

```
plot(hc)
abline(h=8, col = "red", lty =2)
```

Cluster Dendrogram



```
d  
hclust (*, "complete")
```

#Not cut yet, just indicates where we are cutting. Why 8? This is where the two clusters were.

To accomplish this we use `cutree()` function on our `hclust` object.

```
grps <- cutree(hc, h=8)  
grps
```

```
table (grps, km$cluster)
```

grps	1	2
1	30	0
2	0	30

Dimensionality Reduction / Principle Component Analysis (PCA)

- We are attempting to reduce the features of dimensionality while losing as little information as possible.

PCA = Eigenvectors ; Is “like” a big funnel. Where our output allows us to investigate our data.

PCA of UK Food Data

First, import the dataset required.

```
url <- "https://tinyurl.com/UK-foods"  
x <- read.csv(url)  
x
```

		X	England	Wales	Scotland	N.Ireland
1	Cheese		105	103	103	66
2	Carcass_meat		245	227	242	267
3	Other_meat		685	803	750	586
4	Fish		147	160	122	93
5	Fats_and_oils		193	235	184	209
6	Sugars		156	175	147	139
7	Fresh_potatoes		720	874	566	1033
8	Fresh_Veg		253	265	171	143
9	Other_Veg		488	570	418	355
10	Processed_potatoes		198	203	220	187
11	Processed_Veg		360	365	337	334
12	Fresh_fruit		1102	1137	957	674
13	Cereals		1472	1582	1462	1494
14	Beverages		57	73	53	47
15	Soft_drinks		1374	1256	1572	1506
16	Alcoholic_drinks		375	475	458	135
17	Confectionery		54	64	62	41

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

```
dim(x)
```

```
[1] 17 5
```

One potential solution, is to set the row names, doing so “manually”.

```
rownames(x) <- x[,1]
```

Removing the first columns can be accomplished by -1 to start indexing in column.

```
x <- x[,-1]  
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033
Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

The issue with this, you start to lose columns / data.

This can also be accomplished by arguing with `read.csv()`.

```
x <- read.csv(url, row.names = 1)  
x
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139
Fresh_potatoes	720	874	566	1033

Fresh_Veg	253	265	171	143
Other_Veg	488	570	418	355
Processed_potatoes	198	203	220	187
Processed_Veg	360	365	337	334
Fresh_fruit	1102	1137	957	674
Cereals	1472	1582	1462	1494
Beverages	57	73	53	47
Soft_drinks	1374	1256	1572	1506
Alcoholic_drinks	375	475	458	135
Confectionery	54	64	62	41

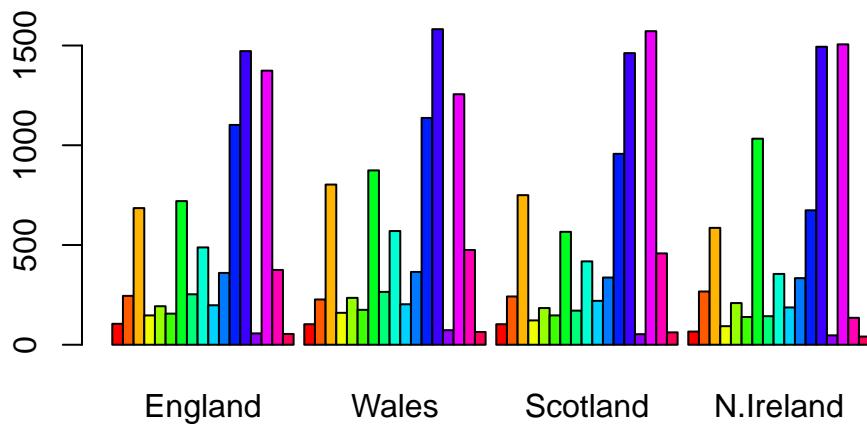
Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

- We rather argue and set the name immediately when we read in the csv. Doing so, allows us to start data d.f manipulations from the get go.

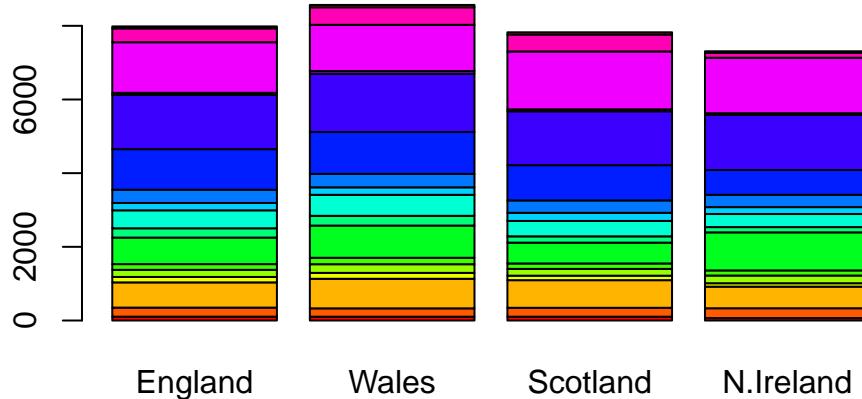
Spotting Major Differences and trends

This is difficult, even in a wee little graph like this with 17D.

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q4: Changing what optional argument in the above ggplot() code results in a stacked barplot figure?

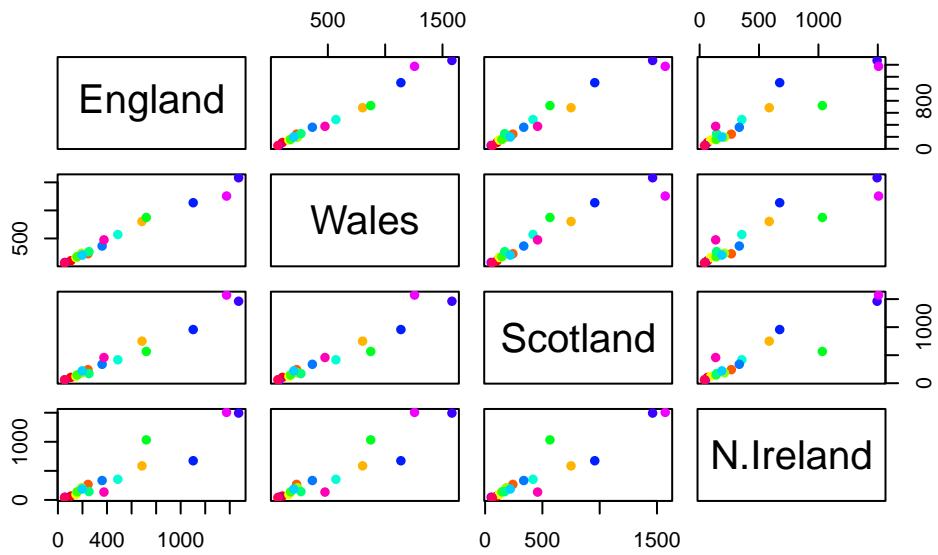
-

Q5: We can use the pairs() function to generate all pairwise plots for our countries. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

-

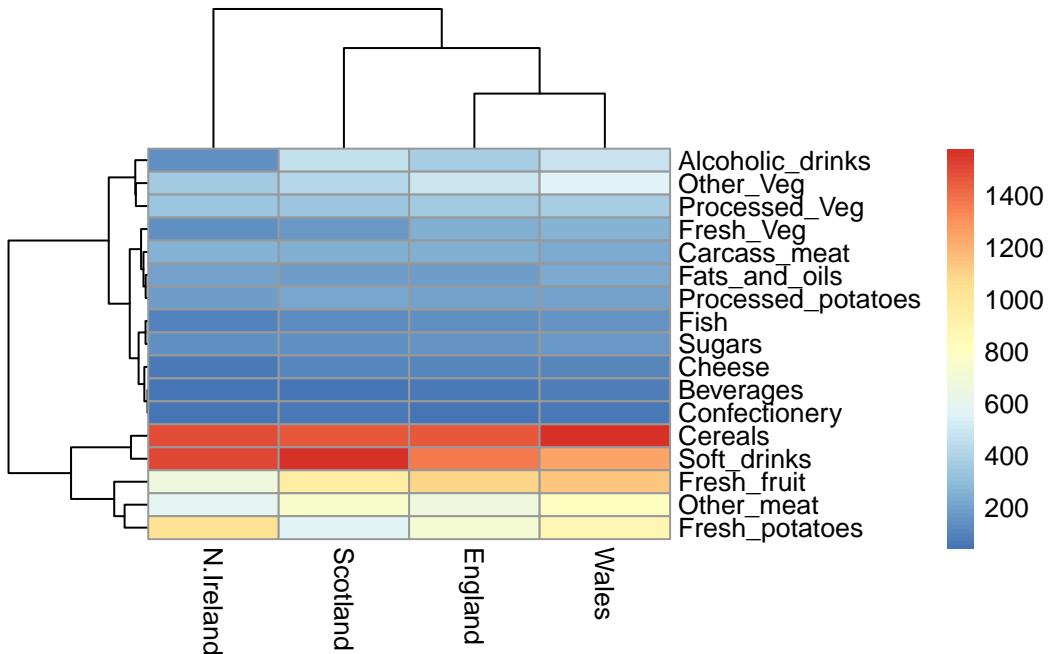
```
pairs(x, col=rainbow(nrow(x)), pch=16)
```

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



Observe the top row, from left to right. On the Y-axis for these plots, you have England's. All dots represent the food categories.

```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Based on the pairs and heatmap figures, which countries cluster together and what does this suggest about their food consumption patterns? Can you easily tell what the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

-

Principal Component Analysis PCA - “To the rescue!”

The main PCA function in **base R**, which is probably the most difficult to use, is ‘prcomp()’. This functions wants the transpose of our food data (i.e. the foods as columns and the counties as rows)

```
pca <- prcomp(t(x)) #t(x) = transpose matrix.
```

```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	3.176e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

```
attributes(pca)

$names
[1] "sdev"      "rotation"   "center"     "scale"      "x"

$class
[1] "prcomp"
```

To make one of our main PCA result figures, we turn to `PCA$x` the scores along our new PCs. This is referred to as score plots, pc plots, ordination plots and various others.

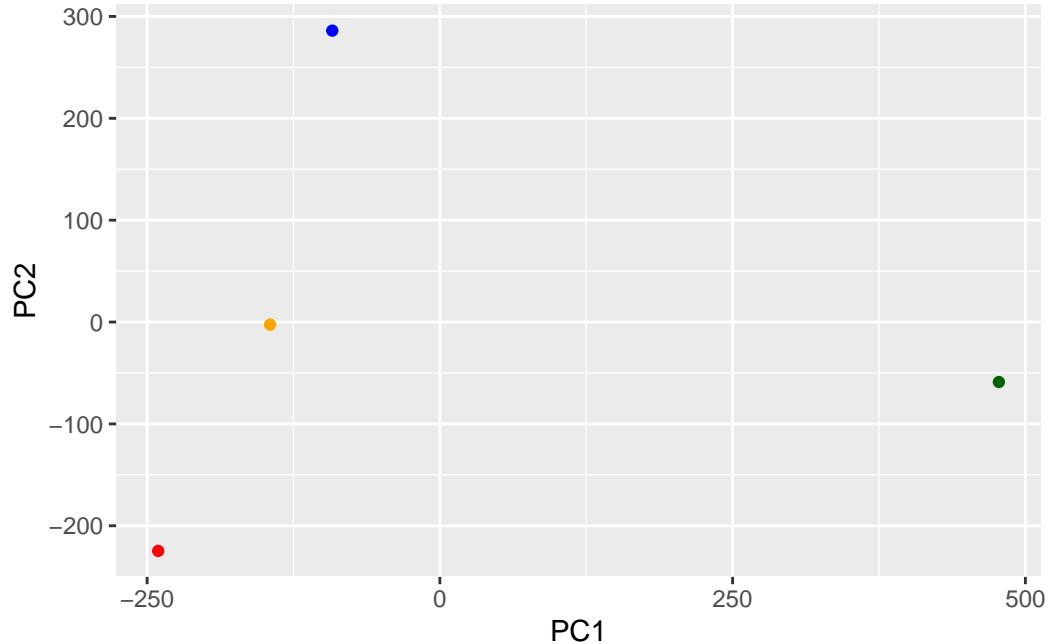
```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-4.894696e-14
Wales	-240.52915	-224.646925	-56.475555	5.700024e-13
Scotland	-91.86934	286.081786	-44.415495	-7.460785e-13
N.Ireland	477.39164	-58.901862	-4.877895	2.321303e-13

```
my_cols <- c("orange", "red", "blue", "darkgreen")

library(ggplot2)

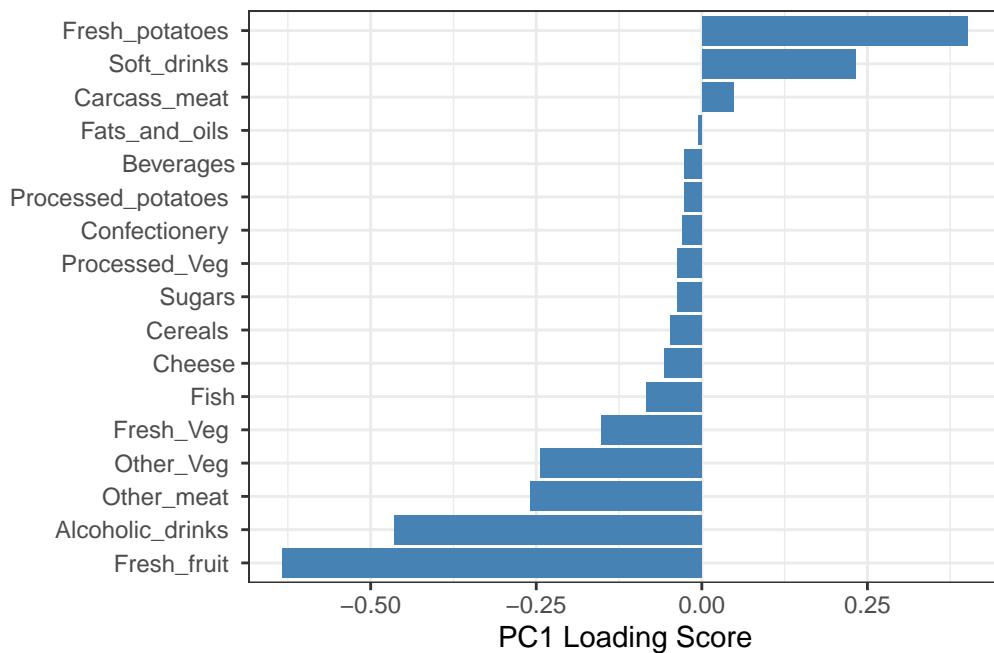
ggplot(pca$x) + aes(PC1, PC2) + geom_point(col=my_cols)
```



The second major result figure is called a “loadings plot” of “variable contributions plot” or “weight plot” # Lots of Terminology

- How do the original variables weigh the respective axis. Think , weighted grading.

```
ggplot(pca$rotation) +
  aes(PC1, y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



This can show us, why are these cells active compared to others, well, due to genes which are up-regulated compared to the other .